

ChildTaskTracker

From working single-family prototype to a networked, multi-household tool — sequenced by what unlocks real use first, and what's genuinely optional.

DRAFT v1.0 2026-09-15 OWNER Adam Helwa COMPANION TO [ChildTaskTracker PRD](#)

● Now — building ● Next — pilot-driven ● Later — only if the pilot demands it



Phase 0 — Prototype

NOW

In progress

Prove the full loop end-to-end on one machine, for one family, at near-zero cost. This is the scope defined in the PRD's MVP section.

SHIPS

- React front end + Node/Express + SQLite, running locally
- Multiple children, mixed-recurrence tasks, points bank, reward catalog
- Child completion → parent validation queue → 24h auto-approval
- Streaks with one-day forgiveness; completion-rate, streak, and points/badge KPIs
- On-demand OpenAI calls for starter tasks, reward recommendations, KPI summaries, reward-frequency tailoring — all constrained to the parent's own reward catalog
- In-app notification center + toast; Parent A / Parent B view switcher
- Age-adaptive theming (playful vs. clean skin)

EXIT CRITERIA

The four usage metrics and four technical metrics in the PRD are all measurable end-to-end in one live session, without a human patching a broken step.



Phase 1 — Family Pilot

NEXT

First 4–6 weeks of real use

Run it with your own kids. This phase produces the data that decides everything after it — no new features unless the pilot surfaces a real gap.

FOCUS

- Track the PRD's usage metrics weekly: completion rate, streak retention, validation turnaround, redemption rate
- Tune AI prompt templates against real KPI summaries and reward recommendations — check they stay inside guardrails and actually sound like something a kid or parent would want to read
- Watch the 24h auto-approve and one-day streak-forgiveness rules against real family rhythm; adjust the thresholds, not the mechanism, if they feel wrong
- Log every corner case actually hit that isn't already in the PRD's corner-case table

EXIT CRITERIA

Either the usage targets are being hit consistently, or you have a specific, named reason they aren't — not a vague "needs polish."



Phase 2 — Real Two-Parent Access

LATER

Only if Phase 1 shows the view-switcher is the bottleneck

Replace the simulated Parent A/B switcher with something that works when parents are actually on separate devices.

FOCUS

- Lightweight parent auth (even just a shared PIN per parent) instead of a view switcher
- Real push notifications (FCM/APNs) replacing the in-app toast as the primary alert
- Backend reachable from both parents' devices — local network first, cheap cloud hosting if the family isn't always on the same WiFi

EXIT CRITERIA

A task marked complete on the child's device produces a real, working notification on both parents' separate phones without anyone opening the app first.



Phase 3 — Richer Insight

LATER

Only once there's enough history to show trends

The KPI set stays deliberately thin in Phase 0–1 (completion rate, streaks, points) because there isn't enough data yet to make more sophisticated views worth the build

cost.

FOCUS

- Category-level KPI breakdowns (e.g. "strong at chores, needs work on homework")
- Multi-week trend charts, once there's more than a few weeks of history to chart
- Revisit AI call cadence — a weekly batch summary may become worth the small added cost once there's a real trend to summarize

EXIT CRITERIA

At least 6–8 weeks of completion history exist for a category breakdown or trend line to say anything a single week's KPIs don't already show.



Phase 4 — Beyond One Household

LATER

Only if there's a reason to extend past this family

Everything here was explicitly deferred in the PRD because it adds real complexity (auth, tenancy, moderation) that a single-family prototype doesn't need.

FOCUS

- Read-only sharing for extended family (grandparents, babysitters)
- Separated-household mode — two independent parent accounts on one shared child record
- Fully custom task scheduling beyond daily/weekday/weekly/one-off
- Multi-tenant architecture, if this ever moves from personal tool toward a product other families could use

EXIT CRITERIA

A specific household or use case actually needs one of these — not speculative demand.

Phase 1 needs

Phase 0 fully working — no pilot without a working loop

Phase 2 needs

Phase 1 evidence that the switcher, not the concept, is the limit

Phase 3 needs

Weeks of real completion data from Phase 1–2 — the feature is worthless without history

Phase 4 needs

A concrete second household or use case — not "eventually might"

Deliberately not on this roadmap

Named so they don't quietly creep back in as "quick additions."

- **App store packaging.** This stays a browser-based, locally-run tool unless a phase explicitly reopens that decision — native builds bring signing, review, and distribution overhead this project isn't scoped for.
- **Monetary or screen-time rewards generated by AI.** The guardrail in the PRD is permanent, not a Phase 0 shortcut — the model recommends only from the parent's own catalog, at every phase.
- **In-app dispute/negotiation threads.** Validation disagreements stay an offline, face-to-face conversation by design, not a feature gap to fill later.